Microsoft

MLOps for R with Azure Machine Learning

David Smith
Cloud Advocate, Microsoft
@revodavid

aka.ms/mlops-r

@revodavid at #rstudioconf

What is DevOps?

People. Process. Products.

> DevOps is the union of **people**, **process**, and **products** to enable continuous delivery of value to your end users.

Donovan Brown, Microsoft
http://bit.ly/WhatIs-DevOps

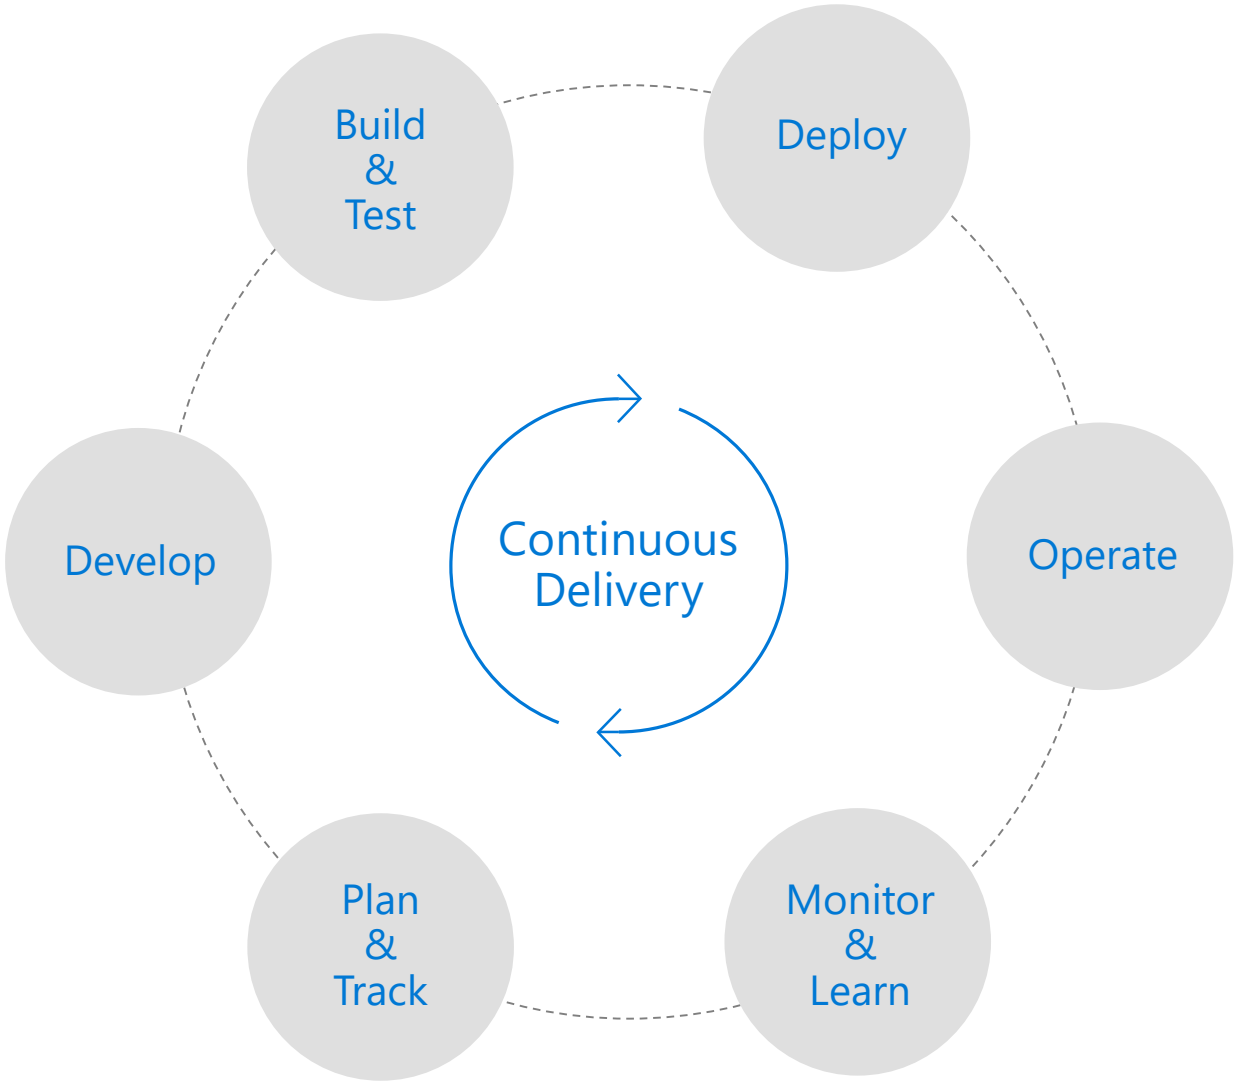

DevOps Process

1 Plan

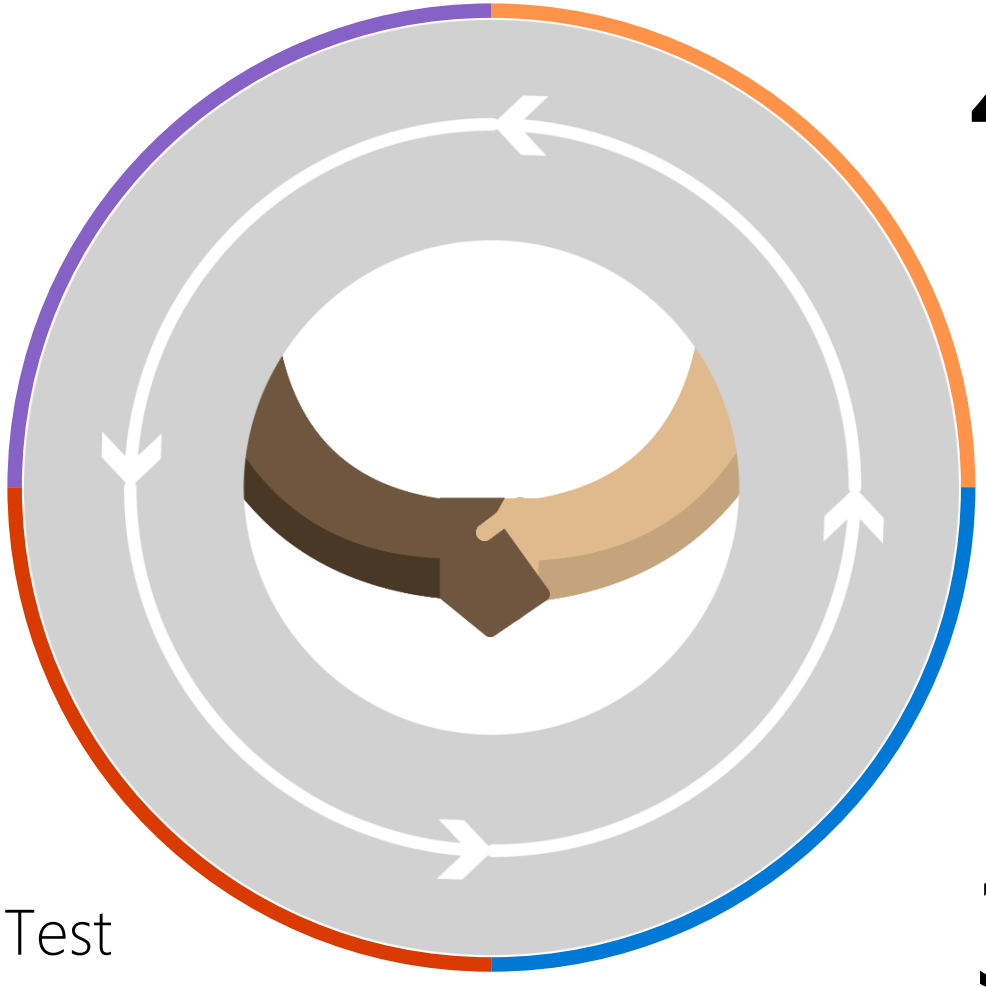

4 Monitor + Learn

Development

Production

2 Develop + Test

3 Release

Machine Learning Process

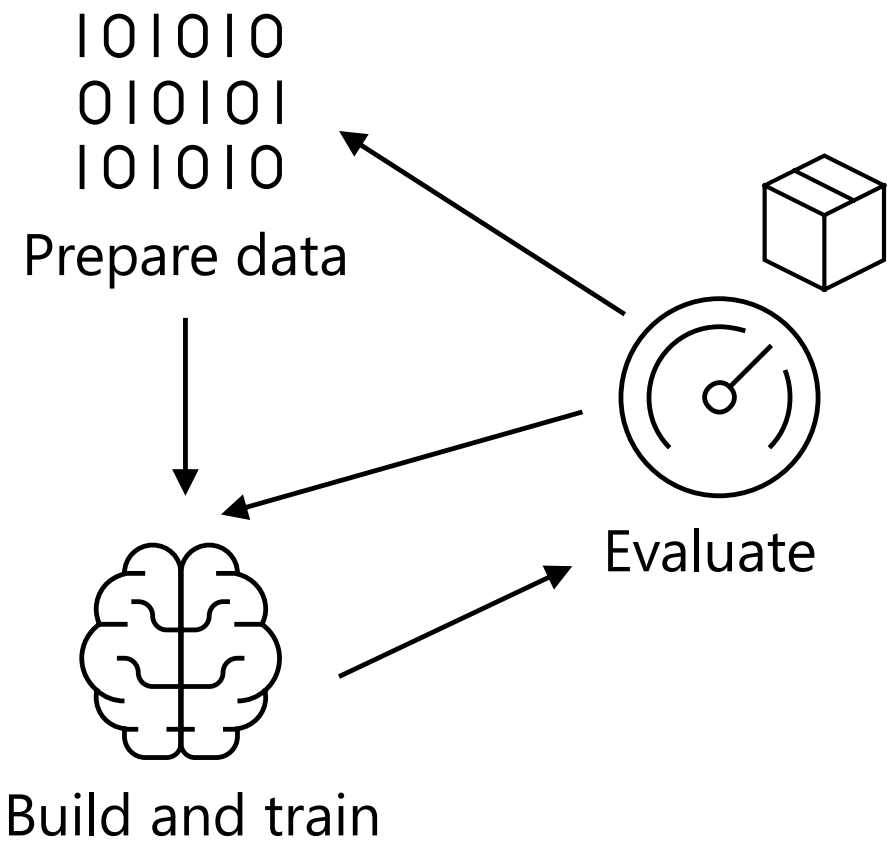

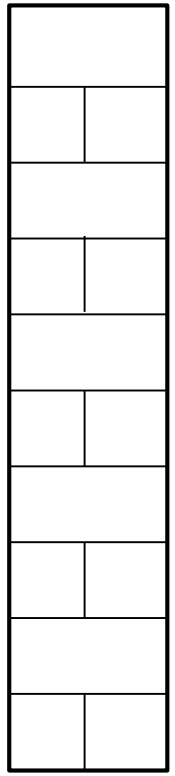

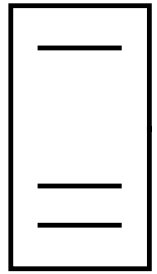

Production

DEVOPS	MLOPS
Manage code (source files)	Manage code (source files) Manage data files, notebooks, Rmd docs
Manage infrastructure (as code)	Manage infrastructure (as code) Manage environments (as code)
Source code control	Source code control Track experiment outcomes Manage data sets
Build executables Builds take hours (mostly) commodity compute	Train models Model training may take weeks or months GPU compute
Manage build versions	Manage model versions Manage reproducible environments
Tests (deterministic) Fix bugs with code	Tests (probabilistic) Fix bugs with code and/or data Model drift / model retraining

Azure Machine Learning Service: ml.azure.com

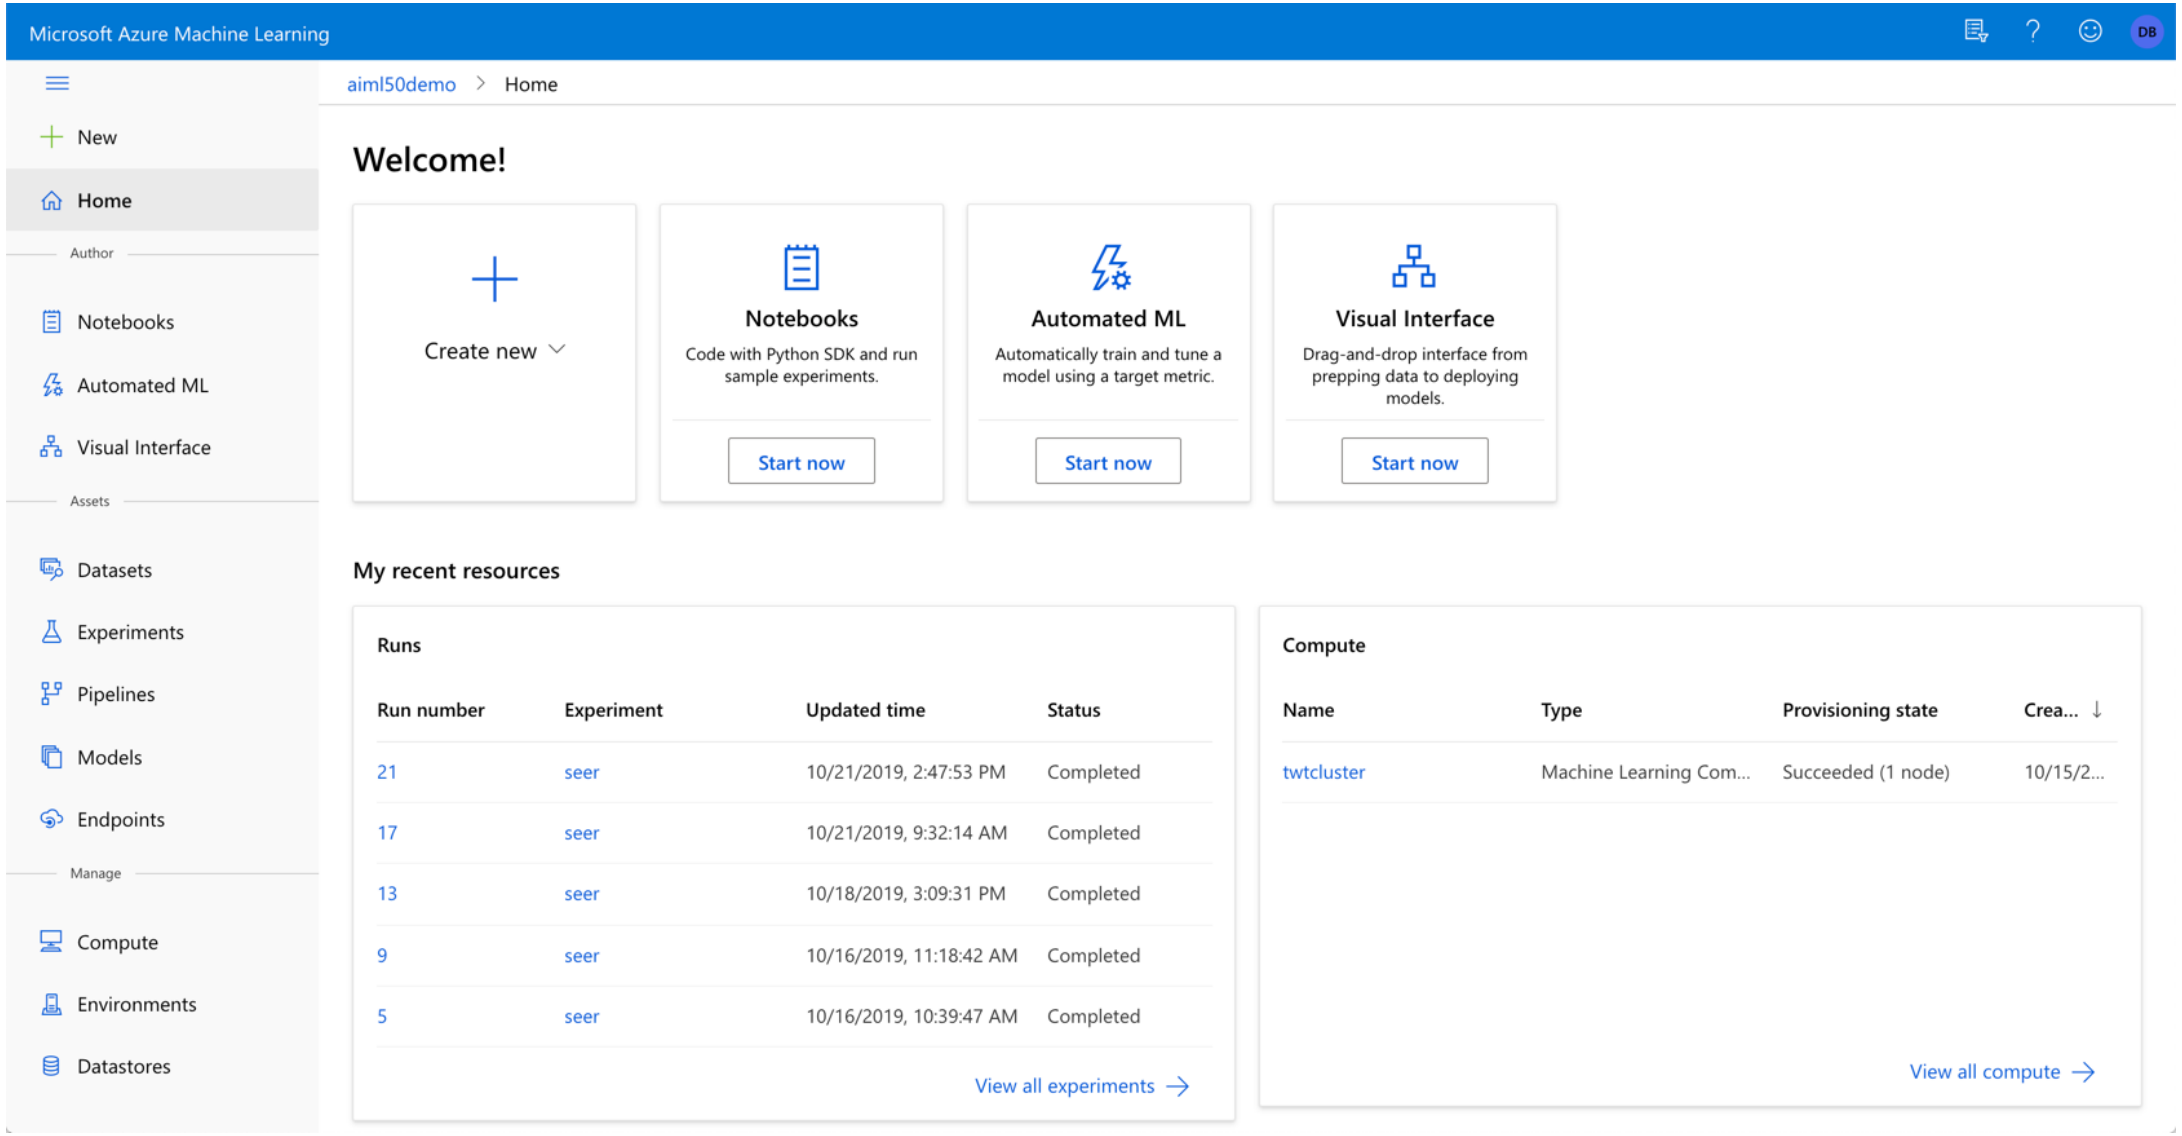

Azure Machine Learning

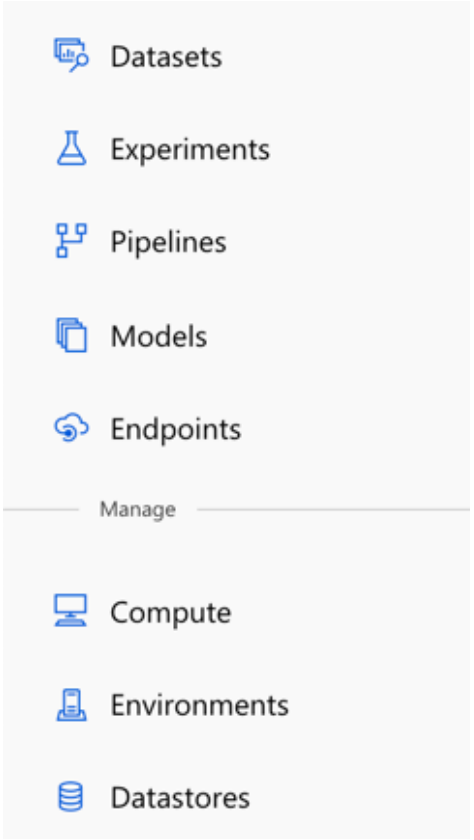

Datasets – registered, known data sets

Experiments – Training runs

Pipelines – Training workflows

Models – Registered, versioned models

Endpoints:

- **Real-time Endpoints** – Deployed model endpoints
- **Pipeline Endpoints** – Training workflow endpoints

Compute – Managed compute

Environments – defined training and inference environments

Datastores – Connections to data

Azure ML service SDK for R

Open-source R package for use with CRAN R: **azuremlsdk**

- Create Workspaces, Experiments, Compute, Models, and other artifacts with R commands
- Use any R function/package (and track requirements for deployment)
- HyperDrive support: smart hyperparameter search with parallel compute
- Publish models as web services (in Azure or your own infra)
- Trigger training / deployment pipelines from CI/CD services

Accident fatality prediction app

Manage Data
- Import flat CSV file and clean data
- Export to datastore

Train Model
- Create training cluster
- Run experiments: GLM, KNN, GLMNET

Deploy model
- Select model by accuracy
- Deploy R function as container

Integrate model
- Shiny application
- Call R function via REST endpoint

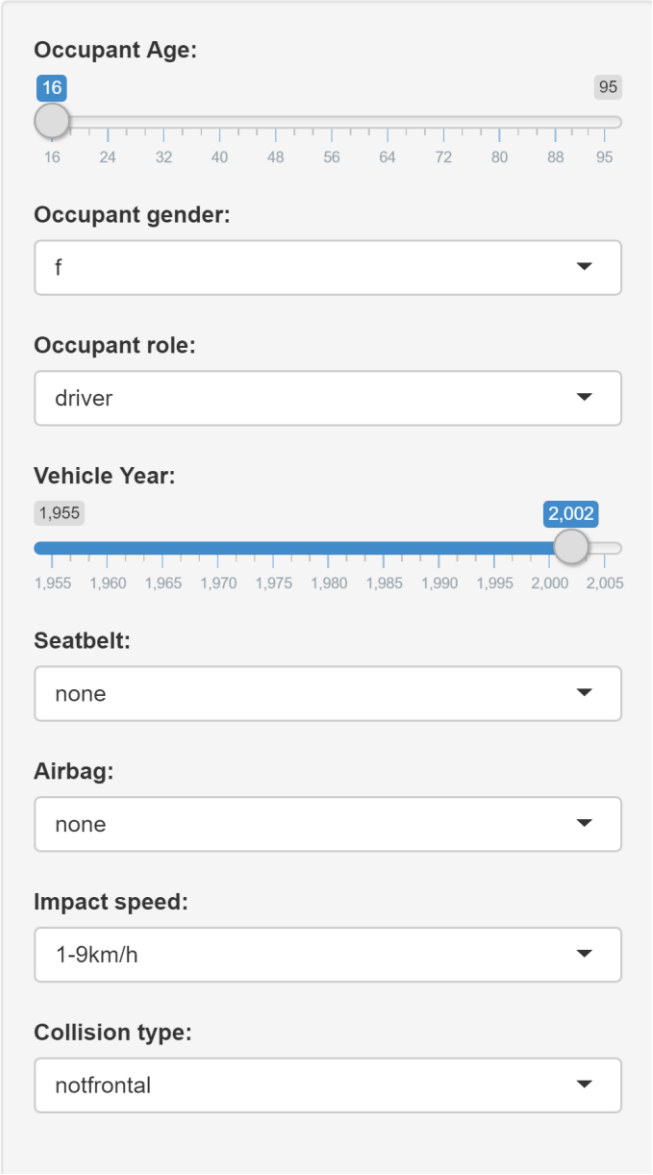

Create compute instance for interactive work

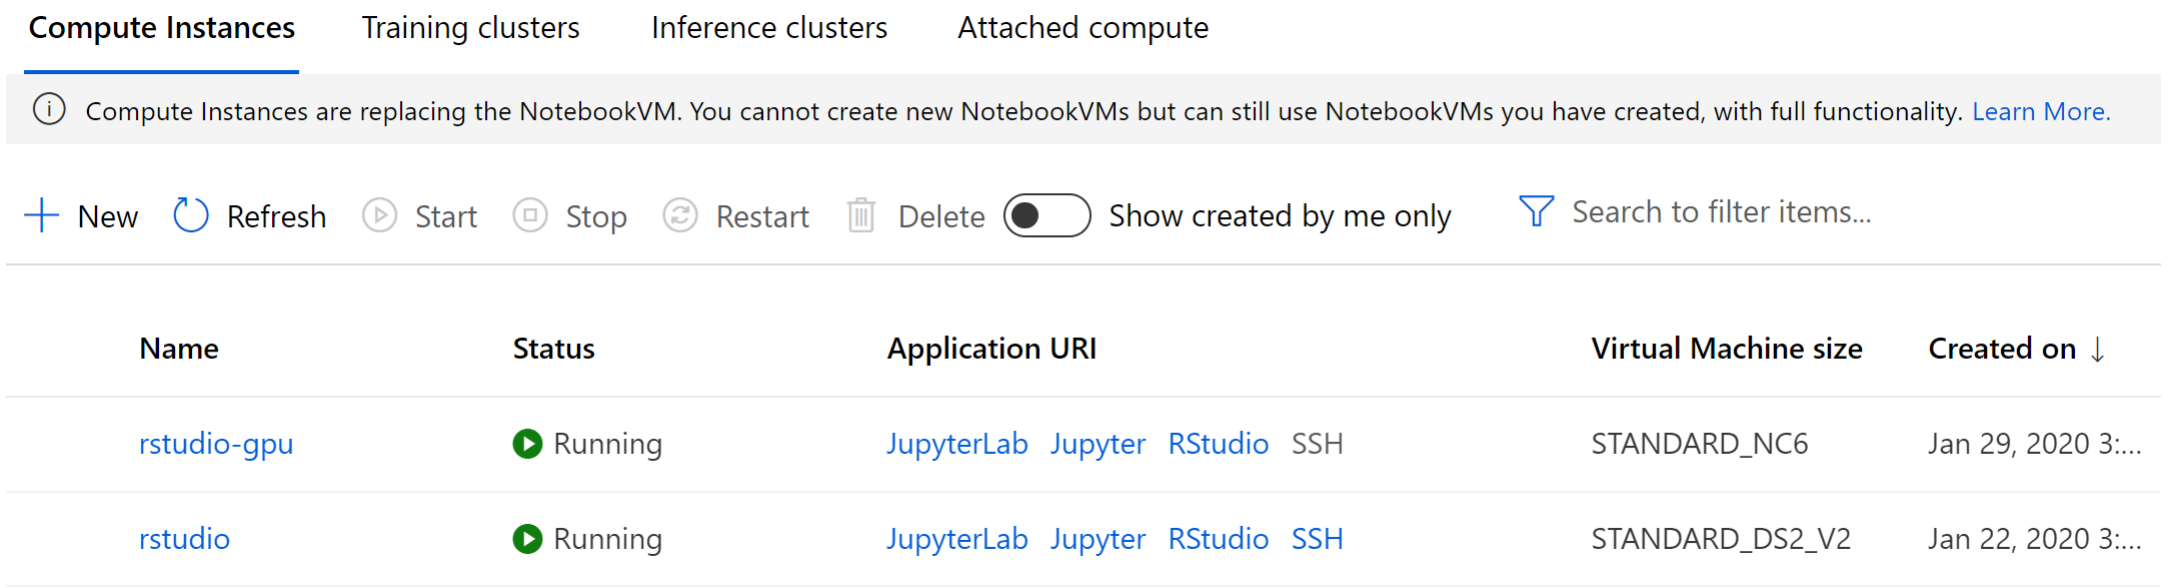

Instances are persistent, charge at usual VM rates

... or just use desktop R/Rstudio with azuremlsdk package

... or use your own attached compute server

Use Compute Instance to prep and share data

```
library(azuremlsdk)
ws <- load_workspace_from_config()
```

azuremlsdk package and auth keys pre-installed on compute instance

```
nassCDS <- read.csv("nassCDS.csv")
# Lots of cleaning code
saveRDS(accidents,
        file="accidents.Rd")
```

Normal R code for data / model experimentation

```
ds <- get_default_datastore(ws)
target_path <- "accidentdata"
upload_files_to_datastore(ds,
          list("./accidents.Rd"),
          target_path = target_path,
          overwrite = TRUE)
```

Export data to shared, persistent datastore
Can use for files other than data, too

Create 2-node training cluster

```
ws <- load_workspace_from_config()
compute_target <- create_aml_compute(
                      workspace = ws,
                      cluster_name = "rcluster",
                      vm_size = "STANDARD_D2_V2",
                      vm_priority = "lowpriority",
                      min_nodes = 1,
                      max_nodes = 2)
```

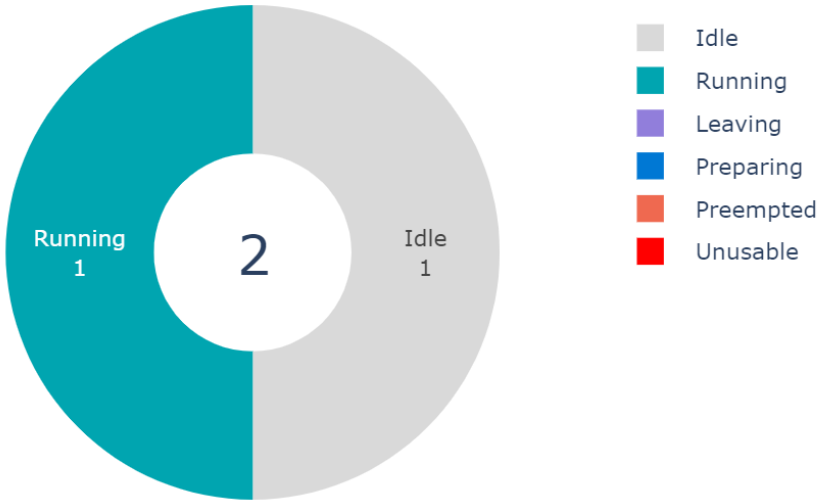

- Nodes pre-loaded with R and R packages
- Scales up as jobs added to queue, scales down when unused
- Resize as needed
- Shared between workspace users (Python and R)
- Low priority nodes for ~80% discount
- GPU clusters available

Train models and choose one to deploy

```
exp <- experiment(ws, "accident")
est <- estimator(source_directory = ".",
                 entry_script = "accident-glmnet.R",
                 script_params = list(
                   "--data_folder" = ds$path(target_path),
                   "--percent_train" = 0.75),
                 compute_target = compute_target)
run.glmnet <- submit_experiment(exp, est)
```

Experiments tracked with source scripts and recorded metrics

Control execution with command line parameters

Most packages pre-loaded, custom packages supported

accident

Switch to old experience

Edit table | Refresh | Reset to default view | Include child runs

Run status

0	30
Running	Completed
3	0
Failed	Other

Method

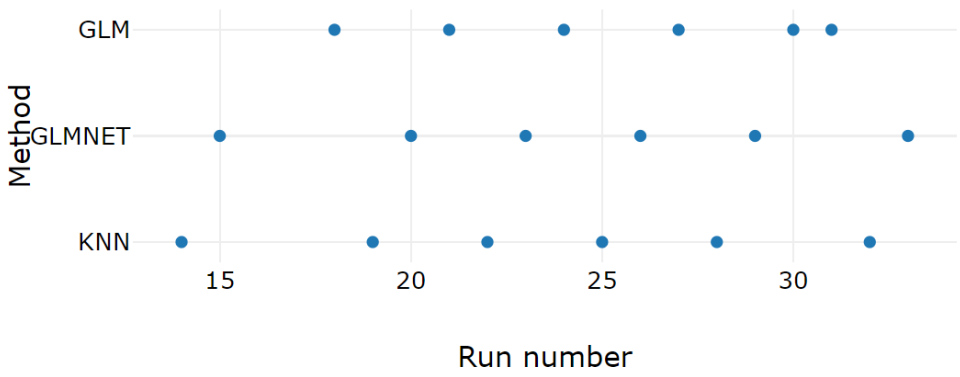

Accuracy

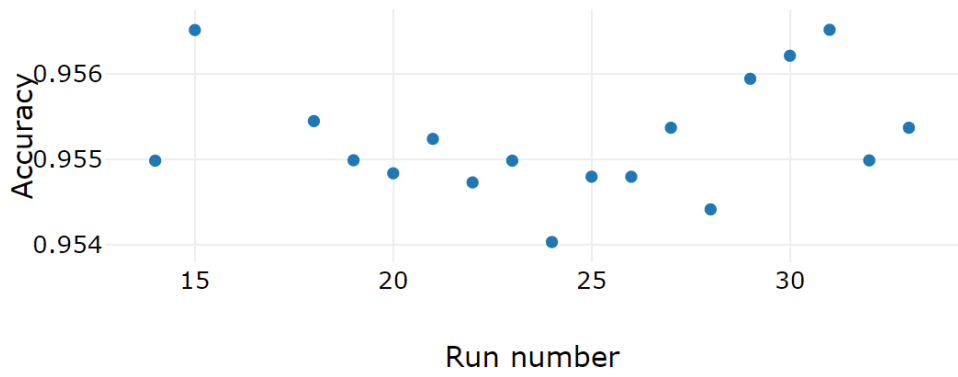

Add filter

Run number	Created time	Duration	Status	Compute target	Run type	Method	Accuracy	TrainPCT
33	Jan 28, 2020 11:39 AM	1m 22s	Completed	rcluster	azureml.scriptrun	GLMNET	0.9553690635132558	0.8
32	Jan 28, 2020 11:39 AM	55s	Completed	rcluster	azureml.scriptrun	KNN	0.9549876025176426	0.8
31	Jan 28, 2020 11:39 AM	54s	Completed	rcluster	azureml.scriptrun	GLM	0.9565134465000954	0.8
30	Jan 28, 2020 11:35 AM	42s	Completed	rcluster	azureml.scriptrun	GLM	0.9562099481232835	0.75

aka.ms/mlops-r

@revodavid at #rstudioconf

Register model

```
model <- register_model(
  ws,
  model_path = "outputs/model.rds",
  model_name = "accidents_model_caret",
  description = "Predict accident probability")
```

Register model with R model object

Models can be versioned

```
r_env <- r_environment(name = "basic_env",
                       cran_packages="caret")

inference_config <- inference_config(
  entry_script = "accident_predict_caret.R",
  source_directory = ".",
  environment = r_env)
```

Packages (CRAN, GitHub or private) needed to execute.

R script receives model and data (JSON) to make prediction

Deploy Model as service

```
aci_config <-
  aci_webservice_deployment_config(
    cpu_cores = 1, memory_gb = 0.5)

aci_service <- deploy_model(ws,
  'accident-pred-caret',
  list(model),
  inference_config,
  aci_config)

accident.endpoint <- get_webservice(
  ws, "accident-pred-caret")$scoring_uri
```

Small container instance for testing

(use Kubernetes for production)

Get REST endpoint for use in app

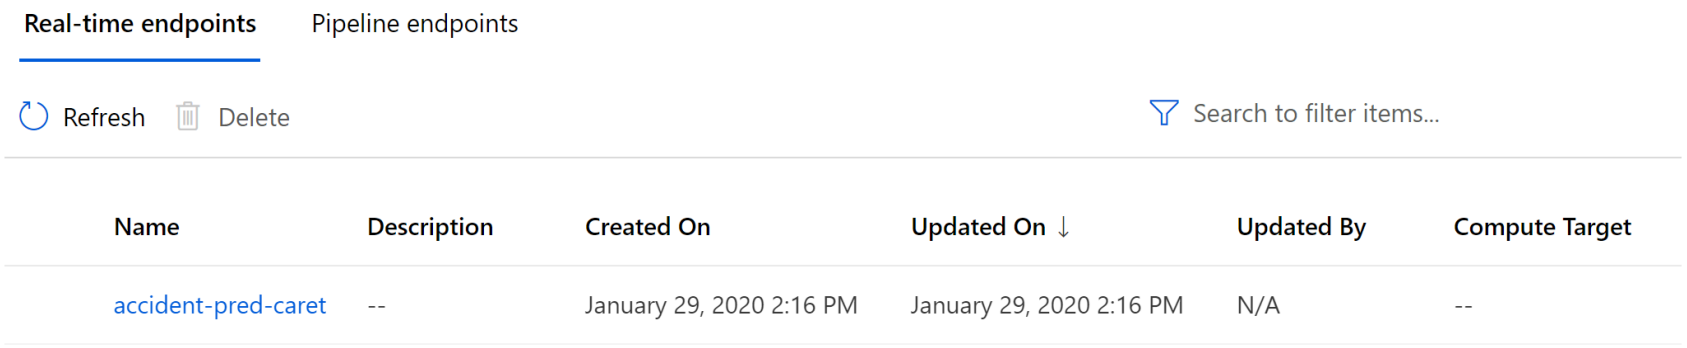

Integrate model into Shiny app

```
library(httr)
v <- POST(accident.endpoint,
          body=input,
          encode="json")

pred <- content(v)[[1]]
```

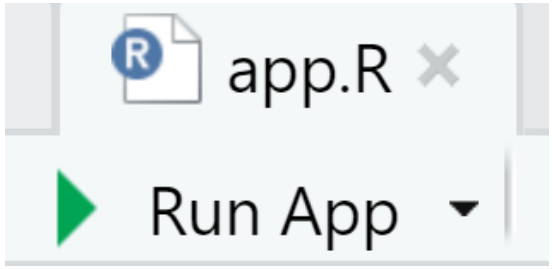

Pass Shiny input data frame to endpoint

Extract prediction from endpoint response

Deploy Shiny app to any server

... or call from any another app

Accident Fatality Probability Estimator

Occupant Age:

16 69 95

16 24 32 40 48 56 64 72 80 88 95

Occupant gender:

f

Occupant role:

pass

Vehicle Year:

1,955 2,002

1,955 1,960 1,965 1,970 1,975 1,980 1,985 1,990 1,995 2,000 2,005

Seatbelt:

belted

Airbag:

none

Impact speed:

40-54

Collision type:

notfrontal

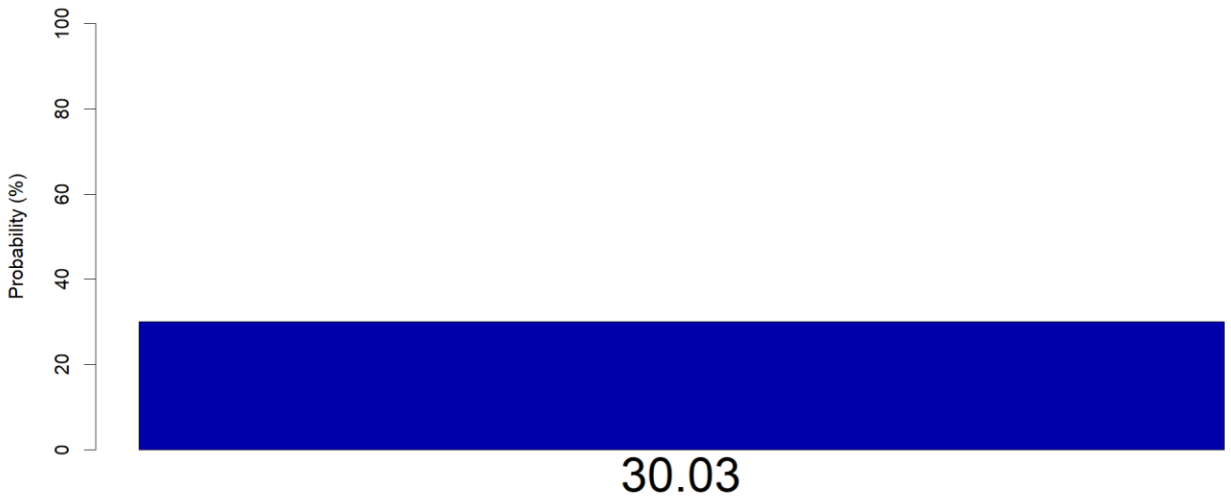

Deliver ML apps with Pipelines

Azure ML Pipelines

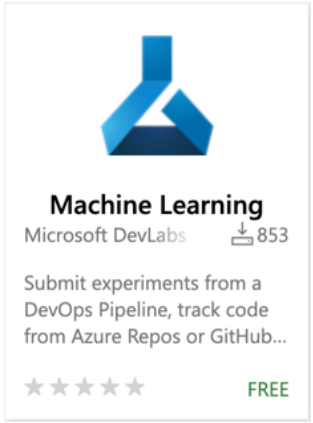


```
az extension add -n azure-cli-ml
```

```
pip install --upgrade azureml-sdk
```

Azure Pipelines
in
Azure DevOps

CI/CD for apps with Azure Pipelines

Azure ML extension triggers builds/releases on model registration in Azure ML Services

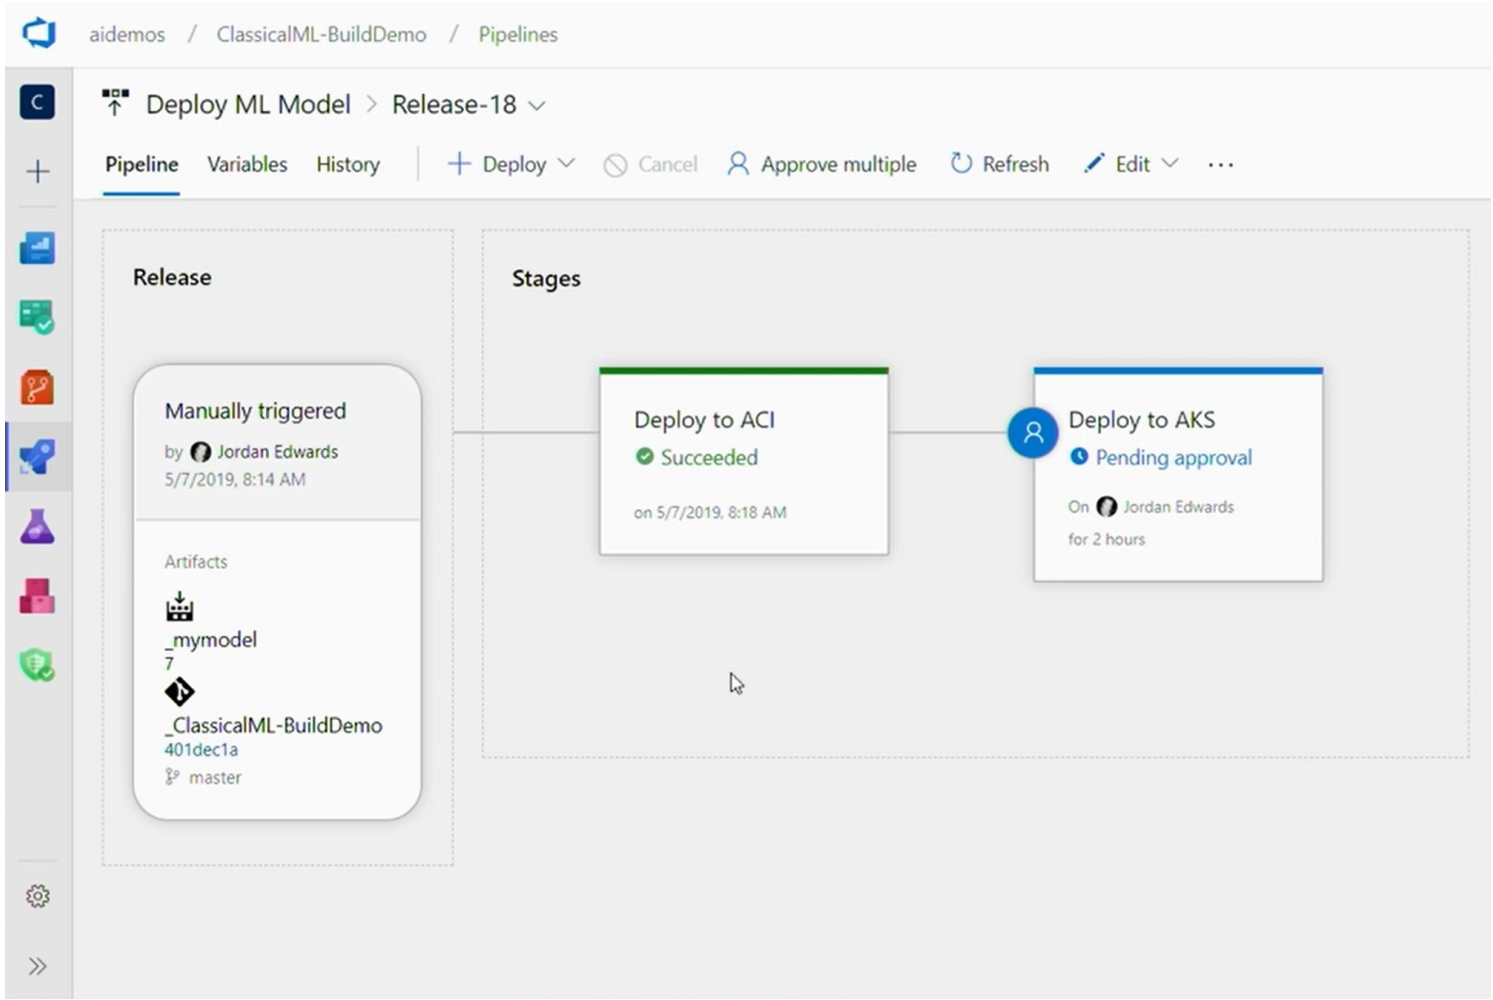

More details:

Machine learning operations: Applying DevOps to data science

aka.ms/AIML50repo

Complete Pipeline

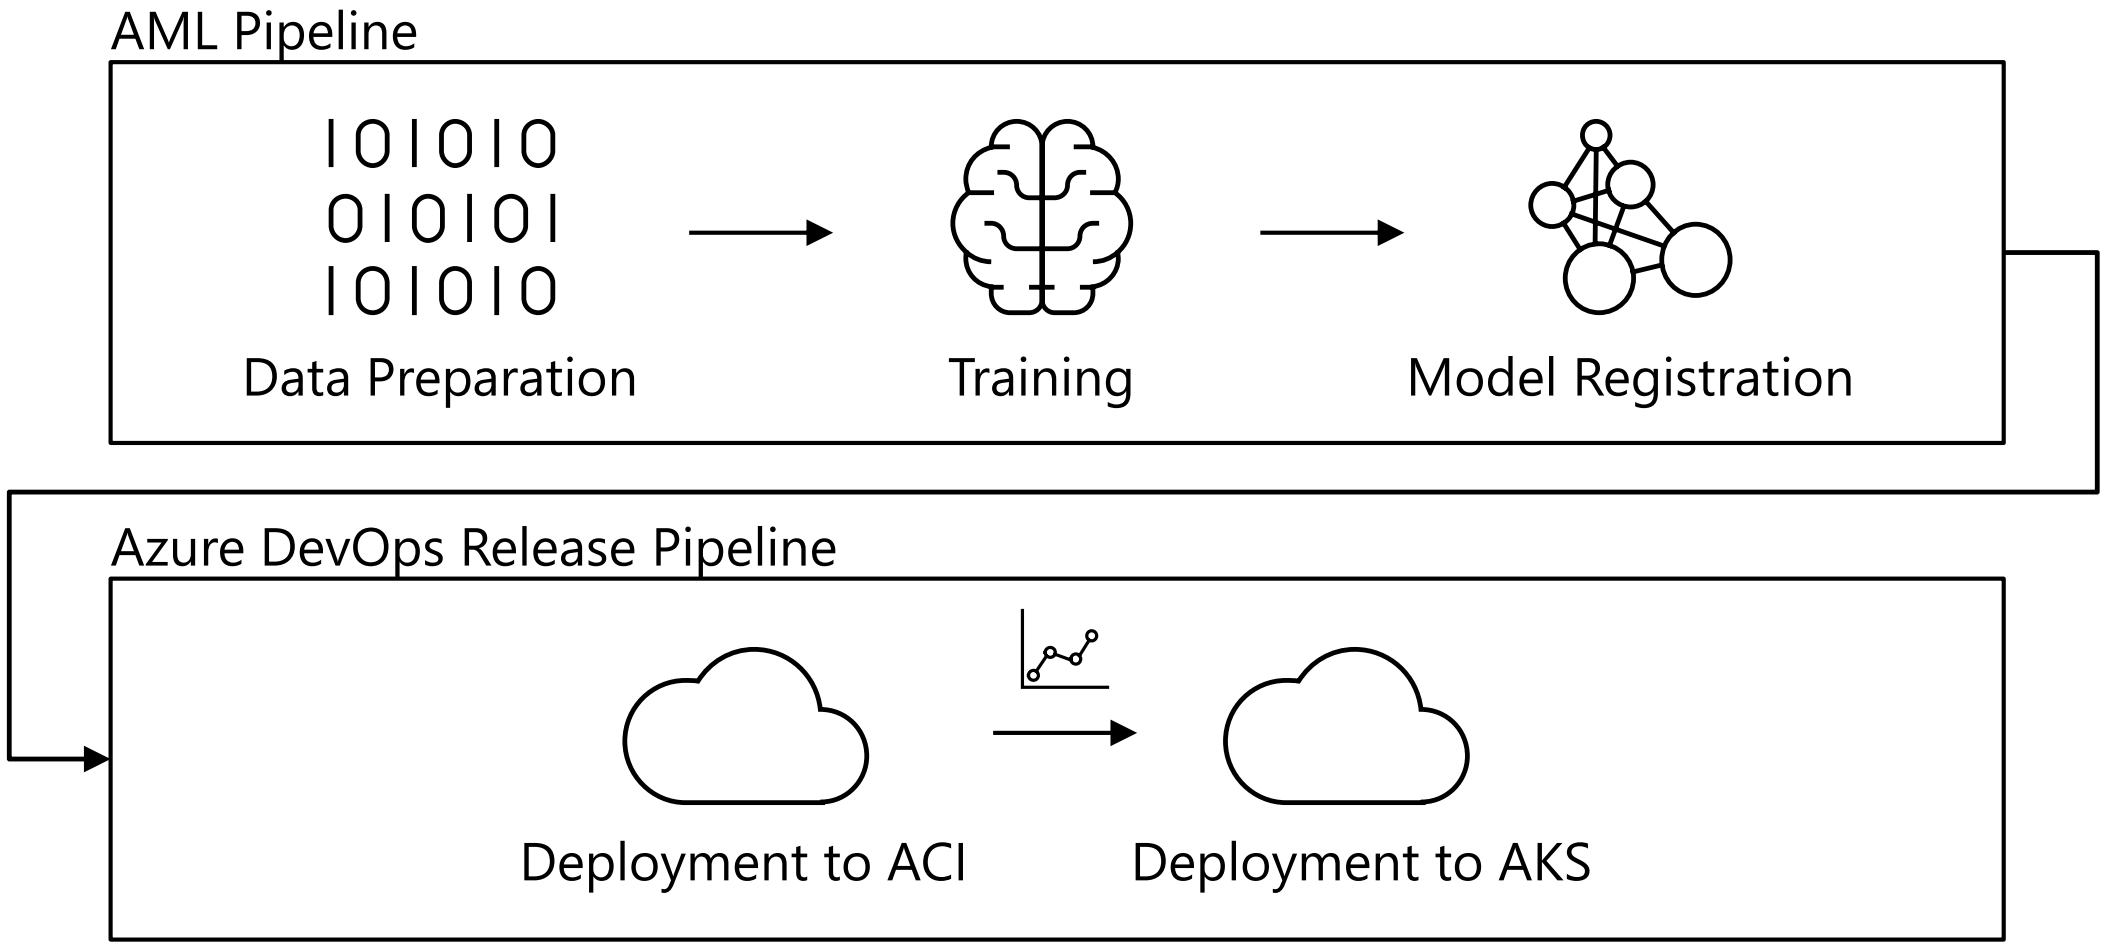

Retraining

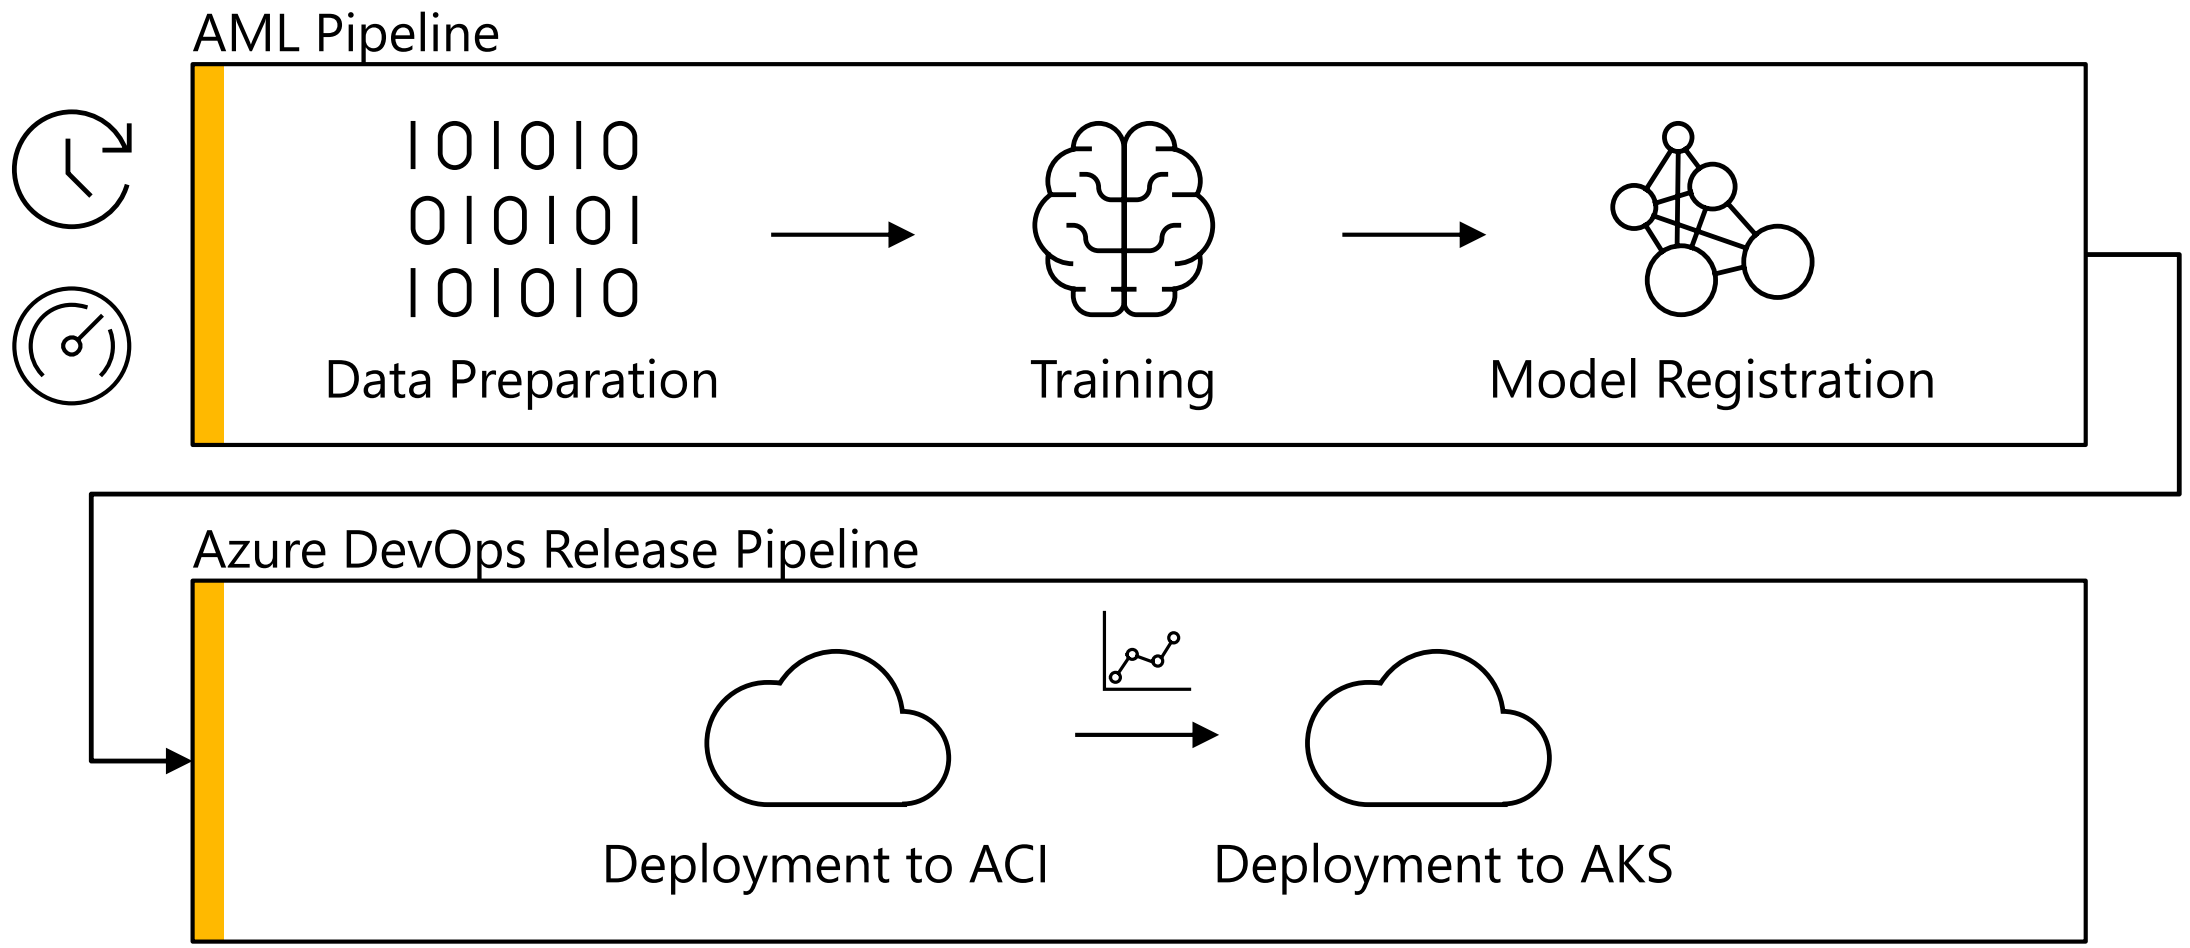
AML Pipeline
Data Preparation
Training
Model Registration
Azure DevOps Release Pipeline
Deployment to ACI
Deployment to AKS

Azure Pipelines

https://azure.com/pipelines

Free **unlimited** build minutes for public projects

Up to 10 free parallel jobs across Windows, Linux and macOS

Azure ML Service

https://azure.com/ml

Free worskpaces, experiments, model registry

Standard Azure rates for compute and deployment **or** use your own servers for free

Thank you!

David Smith
Cloud Advocate, Microsoft
@revodavid

Slides and links:
aka.ms/mlops-r